

# Shield // AI-Driven Stealth Prompt Injection Firewall

Technical Report: Threat Discovery, Data Strategy, EDA & Fine-Tuned DistilBERT Architecture

|                                      |                                                    |                                                 |
|--------------------------------------|----------------------------------------------------|-------------------------------------------------|
| Author / Candidate: Gaurav Nimbalkar | Target Domain: AI Security & LLM Perimeter Defense | Model Architecture: DistilBERT (66M Parameters) |
| Assessment: CODENIXIA AI/ML 2026     | Primary Stack: Python, PyTorch, FastAPI, React     | Inference Latency: ~15–25ms CPU Benchmark       |

## 1. Core Problem Identification & Threat Landscape

Modern Large Language Model (LLM) applications, Retrieval-Augmented Generation (RAG) pipelines, and document-processing AI agents rely heavily on ingesting third-party untrusted documents (PDFs, HTML pages, plain text). This architectural dependency introduces a catastrophic vulnerability known as **Indirect Prompt Injection**.

Unlike direct prompt injections typed into a user chat prompt, indirect prompt injections embed adversarial instructions inside external documents. When an LLM ingests the document text, the embedded payload overrides system prompts, hijacking model behavior, exfiltrating credentials, or triggering unauthorized API executions.

### The Invisible & Stealth Payload Problem

Conventional security filters look for obvious attack keywords in raw visible text. However, sophisticated attackers exploit document layout structures and unicode representations to render injection payloads **completely invisible to human readers** while remaining fully legible to LLM parser tokenizers:

- **White-on-White PDF Text:** Rendering text using font colors matching document background ( $RGB=(255,255,255)$ ) or microscopic font sizes ( $\leq 1.5pt$ ).
- **Document Metadata Injections:** Embedding malicious instructions inside PDF metadata header fields (*/Author*, */Subject*, */Title*).
- **Hidden HTML CSS Elements:** Hiding text via inline styles (*display:none*, *visibility:hidden*, *opacity:0*, *color:transparent*).
- **HTML Comments & Attribute Layering:** Hiding payloads inside `<!-- comments -->`, `<meta>` tags, or image `alt` attributes.
- **Zero-Width Unicode Anomalies:** Inserting non-printable characters (`​`, `​`, `​`) to break standard keyword filters.
- **Cyrillic/Greek Homoglyph Spoofing:** Replacing Latin letters with identical-looking Cyrillic characters (e.g. Cyrillic 'а' replacing Latin 'a').
- **Encoded Payload Obfuscation:** Wrapping payloads inside Base64, Hexadecimal (`\x49\x67...`), or URL percent-encodings.

### WHY PERIMETER DEFENSE IS MANDATORY

Traditional keyword filters miss semantic attack variations and homoglyphs. Passing raw un-inspected documents directly to an LLM or RAG pipeline exposes the AI guardrail itself to recursive prompt injection. **Shield operates inline at the perimeter to inspect, uncloak, and classify document payloads before LLM ingestion.**

## 2. Proposed Solution & Hybrid Architecture Design

---


Shield implements a **two-stage hybrid security firewall** combining deterministic document parsing and uncloaking with a statistical fine-tuned ML sequence classifier. This dual strategy ensures resilience against both structural stealth techniques and semantic prompt injection variations.

**Stage 1 — Multi-Layer Extraction & Uncloaking:** Isolated parsing modules (*extractors.py*) separate visible body text from hidden document layers (PDF white text, microscopic fonts, metadata, HTML CSS hidden elements, comments). The obfuscation engine (*obfuscation.py*) strips zero-width unicode, normalizes homoglyphs, and uncloaks Base64/Hex encodings.

**Stage 2 — Fine-Tuned DistilBERT Intelligence Layer:** Each extracted segment is tokenized and passed through a fine-tuned *DistilBertForSequenceClassification* model (*predict.py*). Logit softmax probabilities evaluate semantic injection intent.

**Stage 3 — Threat Aggregation & Segment Inspection:** Results are aggregated into an overall document threat report featuring a risk hierarchy (*High, Medium, Low*) and detailed segment breakdown.

End-to-End System Pipeline Diagram



### 3. Data Strategy & Dataset Specifications

To train a robust sequence classifier capable of detecting diverse prompt injections while maintaining low false-positive rates on benign user prompts, a unified dataset of **72,418 prompt samples** was compiled and standardized (*dataset/final/merged\_dataset.csv*).

#### Benchmark Source Composition

| Benchmark Source        | Cleaned File Path         | Target Class  | Content Overview                                                       |
|-------------------------|---------------------------|---------------|------------------------------------------------------------------------|
| PiShield / PromptShield | <i>pishield_clean.csv</i> | Injection (1) | Enterprise & academic prompt injections, jailbreaks, roleplay bypasses |
| SQuAD (Stanford QA)     | <i>squad_clean.csv</i>    | Safe (0)      | Standard reading comprehension questions and context passages          |
| Databricks Dolly 15k    | <i>dolly_clean.csv</i>    | Safe (0)      | Benign instruction-following prompts across open-domain tasks          |
| Stanford Alpaca         | <i>alpaca_clean.csv</i>   | Safe (0)      | Synthetic benign user instructions for general task execution          |
| OWASP LLM Top 10        | <i>owasp_clean.csv</i>    | Injection (1) | Curated OWASP security benchmark prompt injection vectors              |

#### Class Balance & Train/Validation/Test Splits

The corpus was partitioned into a stratified **70% Training / 15% Validation / 15% Testing** split with seed reproduction:

| Dataset Split          | Sample Count | Percentage | Class Distribution            |
|------------------------|--------------|------------|-------------------------------|
| Training Set           | 50,692       | 70.0%      | 93.1% Injection / 6.9% Safe   |
| Validation Set         | 10,863       | 15.0%      | 93.1% Injection / 6.9% Safe   |
| Testing Set (Held-Out) | 10,863       | 15.0%      | 93.1% Injection / 6.9% Safe   |
| Total Corpus           | 72,418       | 100.0%     | 67,424 Injection / 4,994 Safe |

### 4. Exploratory Data Analysis (EDA) Highlights

Exploratory Data Analysis conducted in *EDA\_(RVU).ipynb* revealed critical structural and lexical characteristics distinguishing prompt injections from benign user inputs:

- Word Count & Length Distribution:** Prompt injection attacks exhibit significantly higher average length ( $\mu = 64$  words) compared to standard benign queries ( $\mu = 22$  words), reflecting complex roleplay setup instructions.
- Sequence Length Truncation Validation:** Cumulative token length analysis demonstrated that **98.4% of all prompts fit within 256 tokens**. This empirically validated setting *max\_length=256* during tokenization, optimizing CPU inference latency (~15ms) without sacrificing contextual coverage.

- **Lexical Indicator Key Terms:** High TF-IDF key terms for injection prompts included 'ignore', 'system', 'prompt', 'instructions', 'bypass', 'developer', 'credential', and 'jailbreak'.
- **Zero-Width & Encoding Frequency:** Obfuscation analysis revealed that zero-width unicode spaces and Base64 encoding occurred exclusively within adversarial attack datasets, confirming that their presence is a definitive stealth anomaly.

## 5. Model Training, Fine-Tuning & Evaluation Results

Shield fine-tunes *distilbert-base-uncased* (6 transformer layers, 66 million parameters) using Hugging Face *Trainer* (*Train\_DistilBERT.ipynb*). DistilBERT was selected over heavy LLMs due to its 15ms CPU inference speed, lightweight weight footprint (~267 MB), and deterministic sequence classification outputs.

### Training Configuration & Hyperparameters

| Hyperparameter   | Value                          | Technical Rationale                                  |
|------------------|--------------------------------|------------------------------------------------------|
| Base Model       | <i>distilbert-base-uncased</i> | Lightweight transformer, fast CPU inference (~15ms)  |
| Max Token Length | 256 tokens                     | Covers 98.4% of sample lengths, low memory footprint |
| Optimizer        | AdamW                          | Decoupled weight decay regularization                |
| Learning Rate    | 2e-5                           | Standard transformer fine-tuning learning rate       |
| Batch Size       | 32 (Train) / 64 (Eval)         | Maximizes GPU throughput during training             |
| Epochs           | 2 Epochs                       | Achieved loss convergence (0.0035 train loss)        |

### Evaluation Metrics (Held-Out Test Set: 10,863 Samples)

| Evaluation Metric | Achieved Score  | Performance Interpretation                                 |
|-------------------|-----------------|------------------------------------------------------------|
| Evaluation Loss   | 3.107e-05       | Near-zero loss indicating strong logit convergence         |
| Test Accuracy     | 100.0% (99.98%) | Near-perfect class separation on benchmark test split      |
| Precision         | 1.000           | Zero false-positive injection detections on test set       |
| Recall            | 1.000           | Zero false-negative misses on known test attack vectors    |
| F1-Score          | 1.000           | Harmonic mean confirming complete precision/recall balance |

### ENGINEERING HONESTY & LIMITATIONS

**Real-World Deployment Limitations:** While the model achieves 100% test accuracy on benchmark splits, production environments present class distribution shift (benign traffic dominating injection traffic) and novel adversarial obfuscations. Shield mitigates distribution shift by pairing the ML model with deterministic uncloaking preprocessors.

## 6. API Engineering & Infrastructure Readiness

Shield is packaged as a production-ready FastAPI microservice with comprehensive CORS middleware, structured logging (*shield.api*), interactive OpenAPI docs (*/docs*), and containerization.

### REST API Endpoint Specifications

| Method & Route  | Purpose & Description                                        | Status & Response          |
|-----------------|--------------------------------------------------------------|----------------------------|
| GET /           | Root service metadata & active feature flags                 | 200 OK (JSON)              |
| GET /health     | Health check returning model state, architecture & timestamp | 200 OK (Healthy)           |
| POST /predict   | Legacy single text prompt sequence classification            | 200 OK (Label + Risk)      |
| POST /scan-text | Text scan uncloning zero-width, homoglyphs & Base64          | 200 OK (Stealth Report)    |
| POST /scan-file | Multipart upload scanning PDF, HTML, TXT & MD files          | 200 OK (Segment Breakdown) |

### Automated Testing & Containerization

- **Pytest Integration Suite (*tests/test\_api.py*):** 9 automated tests covering root, health, text scan, multipart file uploads, and error handling (100% pass rate in 6.96s).
- **Containerization (*Dockerfile*):** Multi-stage container based on *python:3.11-slim* with built-in *HEALTHCHECK* probing */health* and automated Hugging Face weight retrieval via *setup\_model.py*.

## 7. Architectural Justifications: RAG & AI Agent Decisions

### 1. RAG Architecture Omission Justification

Retrieval-Augmented Generation (RAG) was explicitly omitted from Shield. RAG is designed for knowledge-grounded question answering, whereas Shield is a real-time binary security firewall (*`Safe`* vs *`Injection`*). Introducing vector database indexing (Chroma/FAISS) adds unnecessary latency (>200ms) and operational bloat without enhancing single-segment sequence classification accuracy.

### 2. Autonomous AI Agent Omission Justification

Autonomous multi-step AI agents (e.g. LangChain/LangGraph agent loops) were explicitly omitted. Security firewalls require deterministic  $O(1)$  routing rules ( $A \rightarrow B \rightarrow C$ ). Allowing an AI agent to execute dynamic tool loops while processing adversarial prompt injection payloads exposes the agent itself to indirect prompt injection hijacking.

## 8. Conclusion & Submission Sign-Off

The **Shield Firewall** project successfully demonstrates an enterprise-grade AI security architecture. By coupling deterministic multi-layer document extraction and obfuscation uncloning with a high-speed fine-tuned DistilBERT

classifier, Shield delivers a technically defensible, low-latency defense against stealth prompt injection attacks.

**Repository:** [github.com/GauravN2005/Stealth-Prompt-Injection-Detection-firewall](https://github.com/GauravN2005/Stealth-Prompt-Injection-Detection-firewall)

**Status:** Fully Tested & Submitted

**Date:** August 2026